



# containerd Internals:


## Building a Core Container Runtime

Stephen Day (@stevvooe), Docker

October 24, 2017

Docker Prague Meetup

# A Brief History



APRIL 2016 ..... Containerd "0.2" announced, Docker 1.11

***Management/Supervisor for the OCI `runC` executor***

Announce expansion of containerd OSS project ..... DECEMBER 2016



***Containerd 1.0: A core container runtime project for the industry***

MARCH 2017 ..... Containerd project contributed to CNCF



<https://github.com/containerd/containerd>

containerd / containerd

Unwatch

167

Unstar

1,800

Fork

378

Code

Issues 84

Pull requests 16

Projects 0

Wiki

Insights

An open and reliable container runtime <https://containerd.io>

containerd

oci

containers

docker

cncf

2,673 commits

6 branches

25 releases

104 contributors

Apache-2.0

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

mlaventure Merge pull request #1665 from crosbymichael/bump-runc

Latest commit 3679a55 3 days ago

api

Refactor differ into separate package

12 days ago

archive

Merge pull request #1631 from dmcgowan/cancel-unpack

6 days ago

cmd

Merge pull request #1652 from crosbymichael/cr-image

5 days ago


# Why Containerd 1.0?

- Continue projects **spun out** from monolithic Docker engine
- Expected use **beyond** Docker engine (Kubernetes CRI)
- Donation to **foundation** for broad industry collaboration
  - Similar to runc/libcontainer and the OCI



# Technical Goals/Intentions

- Clean **gRPC-based** API + client library
- Full **OCI** support (runtime and image spec)
- **Stability** and **performance** with tight, well-defined core of container function
- **Decoupled** systems (image, filesystem, runtime) for pluggability, reuse



# Requirements

- **A la carte**: use only what is required
- Runtime **agility**: fits into different platforms
  - Pass-through container configuration (direct OCI)
- **Decoupled**
- Use **known-good** technology
  - OCI container runtime and images
  - gRPC for API
  - Prometheus for Metrics





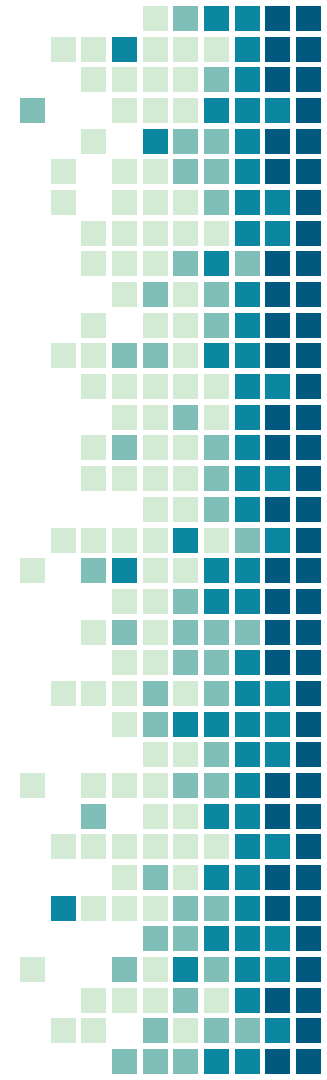
## Use cases

### - **CURRENT**

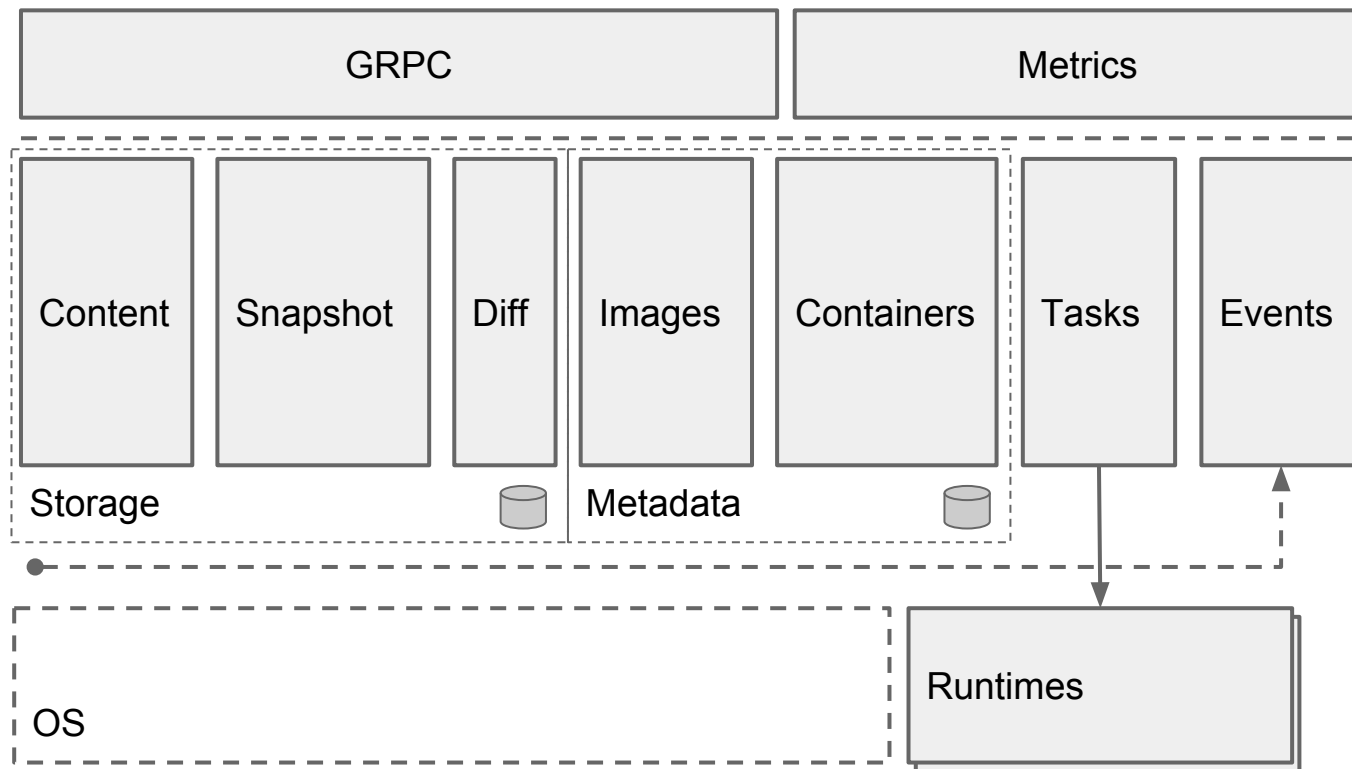
- Docker (moby)
- Kubernetes (cri-containerd)
- SwarmKit (experimental)
- LinuxKit
- BuildKit

### - **FUTURE/POTENTIAL**

- IBM Cloud/Bluemix
- OpenFaaS
- Puppet R&D
- *{your project here}*




# Architecture





# Architecture



```
# HELP container_blkio_io_service_bytes_recursive_bytes The blkio io service bytes recursive
# TYPE container_blkio_io_service_bytes_recursive_bytes gauge
container_blkio_io_service_bytes_recursive_bytes{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Async"} 1.07159552e+08
container_blkio_io_service_bytes_recursive_bytes{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Read"} 0
container_blkio_io_service_bytes_recursive_bytes{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Sync"} 81920
container_blkio_io_service_bytes_recursive_bytes{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Total"} 1.07241472e+08
container_blkio_io_service_bytes_recursive_bytes{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Write"} 1.07241472e+08
# HELP container_blkio_io_serviced_recursive_total The blkio io serviced recursive
# TYPE container_blkio_io_serviced_recursive_total gauge
container_blkio_io_serviced_recursive_total{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Async"} 892
container_blkio_io_serviced_recursive_total{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Read"} 0
container_blkio_io_serviced_recursive_total{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Sync"} 888
container_blkio_io_serviced_recursive_total{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Total"} 1780
container_blkio_io_serviced_recursive_total{container_id="foo4",device="/dev/nvme0n1",major="259",minor="0",namespace="default",op="Write"} 1780
# HELP container_cpu_kernel_nanoseconds The total kernel cpu time
# TYPE container_cpu_kernel_nanoseconds gauge
container_cpu_kernel_nanoseconds{container_id="foo4",namespace="default"} 2.6e+08
# HELP container_cpu_throttle_periods_total The total cpu throttle periods
# TYPE container_cpu_throttle_periods_total gauge
container_cpu_throttle_periods_total{container_id="foo4",namespace="default"} 0
# HELP container_cpu_throttled_periods_total The total cpu throttled periods
# TYPE container_cpu_throttled_periods_total gauge
container_cpu_throttled_periods_total{container_id="foo4",namespace="default"} 0
# HELP container_cpu_throttled_time_nanoseconds The total cpu throttled time
# TYPE container_cpu_throttled_time_nanoseconds gauge
container_cpu_throttled_time_nanoseconds{container_id="foo4",namespace="default"} 0
# HELP container_cpu_total_nanoseconds The total cpu time
# TYPE container_cpu_total_nanoseconds gauge
container_cpu_total_nanoseconds{container_id="foo4",namespace="default"} 1.003301578e+09
# HELP container_cpu_user_nanoseconds The total user cpu time
# TYPE container_cpu_user_nanoseconds gauge
container_cpu_user_nanoseconds{container_id="foo4",namespace="default"} 7e+08
# HELP container_hugetlb_failcnt_total The hugetlb failcnt
# TYPE container_hugetlb_failcnt_total gauge
container_hugetlb_failcnt_total{container_id="foo4",namespace="default",page="1GB"} 0
container_hugetlb_failcnt_total{container_id="foo4",namespace="default",page="2MB"} 0
# HELP container_hugetlb_max_bytes The hugetlb maximum usage
# TYPE container_hugetlb_max_bytes gauge
container_hugetlb_max_bytes{container_id="foo4",namespace="default",page="1GB"} 0
container_hugetlb_max_bytes{container_id="foo4",namespace="default",page="2MB"} 0
# HELP container_hugetlb_usage_bytes The hugetlb usage
# TYPE container_hugetlb_usage_bytes gauge
container_hugetlb_usage_bytes{container_id="foo4",namespace="default",page="1GB"} 0
container_hugetlb_usage_bytes{container_id="foo4",namespace="default",page="2MB"} 0
# HELP container_memory_active_anon_bytes The active_anon amount
# TYPE container_memory_active_anon_bytes gauge
container_memory_active_anon_bytes{container_id="foo4",namespace="default"} 2.658304e+06
# HELP container_memory_active_file_bytes The active_file amount
# TYPE container_memory_active_file_bytes gauge
container_memory_active_file_bytes{container_id="foo4",namespace="default"} 7.319552e+06
# HELP container_memory_cache_bytes The cache amount used
# TYPE container_memory_cache_bytes gauge
container_memory_cache_bytes{container_id="foo4",namespace="default"} 5.0597888e+07
# HELP container_memory_dirty_bytes The dirty amount
```

# Metrics



# Containerd: Rich Go API

## **Getting Started**

<https://github.com/containerd/containerd/blob/master/docs/getting-started.md>

## **GoDoc**

<https://godoc.org/github.com/containerd/containerd>



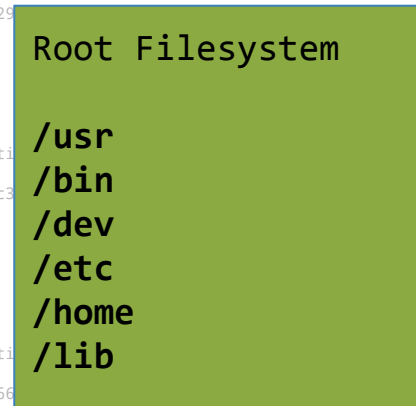
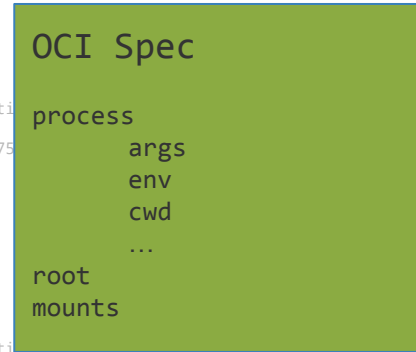


# Pulling an Image

What do runtimes need?



# Image Formats

## Docker and OCI



```
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.list.v2+json",
  "manifests": [
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v1+json",
      "size": 2094,
      "digest": "sha256:7820f9a86d4ad15a2c4f0c0e5479298df2aa7c2f6871288e2ef8546f3e7b6783",
      "platform": {
        "architecture": "ppc64le",
        "os": "linux"
      }
    },
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v1+json",
      "size": 2094,
      "digest": "sha256:ae1b0e06e8ade3a11267564a26e75...",
      "platform": {
        "architecture": "amd64",
        "os": "linux"
      }
    },
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v1+json",
      "size": 2094,
      "digest": "sha256:e4c0df75810b953...",
      "platform": {
        "architecture": "arm64",
        "os": "linux"
      }
    },
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v1+json",
      "size": 2094,
      "digest": "sha256:07ebe243465ef4a...",
      "platform": {
        "architecture": "arm",
        "os": "linux",
        "variant": "armv7"
      }
    },
    {
      "mediaType": "application/vnd.docker.distribution.manifest.v1+json",
      "size": 2090,
      "digest": "sha256:fb2fc0707b86dafa9959fe3d29e66...",
      "platform": {
        "architecture": "arm64",
        "os": "linux",
        "variant": "armv8"
      }
    }
  ]
}
```

# Content Addressability

digest.FromString("foo") ->

"sha256:2c26b46b68ffc68ff99b453c1d30413413422d706483bfa0f98a5e886266e7ae"

digest.FromString("foo tampered") ->

"sha256:51f7f1d1f6bebed72b936c8ea257896cb221b91d303c5b5c44073fce33ab8dd8"

digest.FromString("bar sha256:2c...") ->


"sha256:2e94890c66fbcccca9ad680e1b1c933cc323a5b4bcb14cc8a4bc78bb88d41055"

"bar sha256:2c..."




"foo"

"bar sha256:2c..."

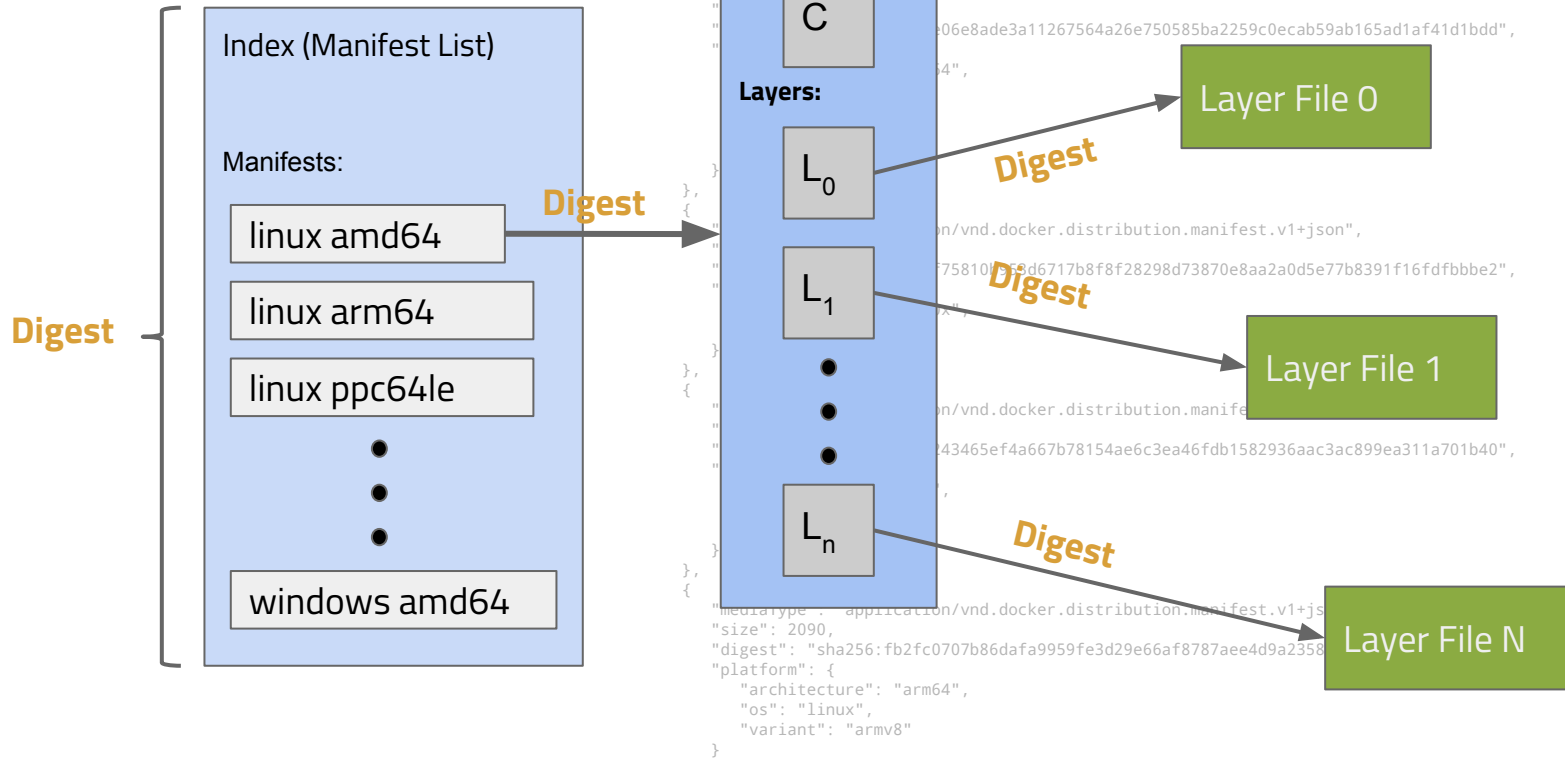


"foo tampered"



# Image Formats


## Docker and OCI



# Resolution

Getting a digest from a name:

**ubuntu**



**sha256:71cd81252a3563a03ad8daee81047b62ab5d892ebbfbf71cf53415f29c130950**



# Image Names in Docker

Reference Type	CLI	Canonical
Repository	ubuntu	docker.io/library/ubuntu
Untagged	ubuntu	docker.io/library/ubuntu:latest
Tagged	ubuntu:16.04	docker.io/library/ubuntu:16.04
Content Trust	ubuntu:latest	docker.io/library/ubuntu@sha256:...
By digest	ubuntu@sha256:....	docker.io/library/ubuntu@sha256:...
Unofficial tagged	stevvooe/ubuntu:latest	docker.io/stevvooe/ubuntu:latest
Private registry tagged	myregistry.com/repo:latest	myregistry.com/repo:latest

# Other Approaches to Naming

- **Self Describing**
  - Massive collisions
  - Complex trust scenarios
- **URI Schemes: `docker://docker.io/library/ubuntu`**
  - Redundant
  - Confuses protocols and formats
  - Operationally Limiting
    - let configuration choose protocol *and* format



# Locators

## Schema-less URIs

ubuntu (Docker name)

↳ docker.io/library/ubuntu:latest (Docker canonical)

↳ (docker.io/library/ubuntu, latest)

locator object



# Remotes *(hint: think git remotes)*


## Locators and Resolution

```
fetcher := resolver.Resolve("docker.io/library/ubuntu")
```

Endlessly Configurable!


```
type Resolver interface {  
    Resolve(ctx context.Context, locator string) (Fetcher, error)  
}
```

```
type Fetcher interface {  
    Fetch(ctx context.Context, id string, hints ...string) (io.ReadCloser, error)  
}
```




# Pulling an Image

## Data Flow

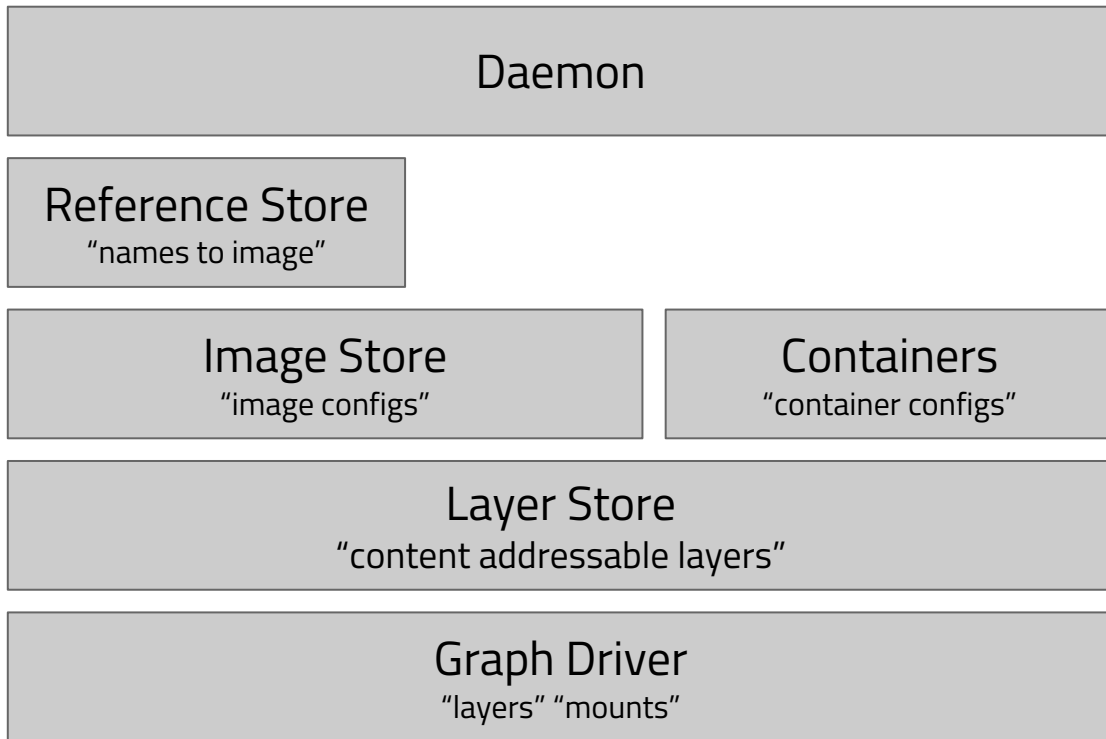


# Snapshotters

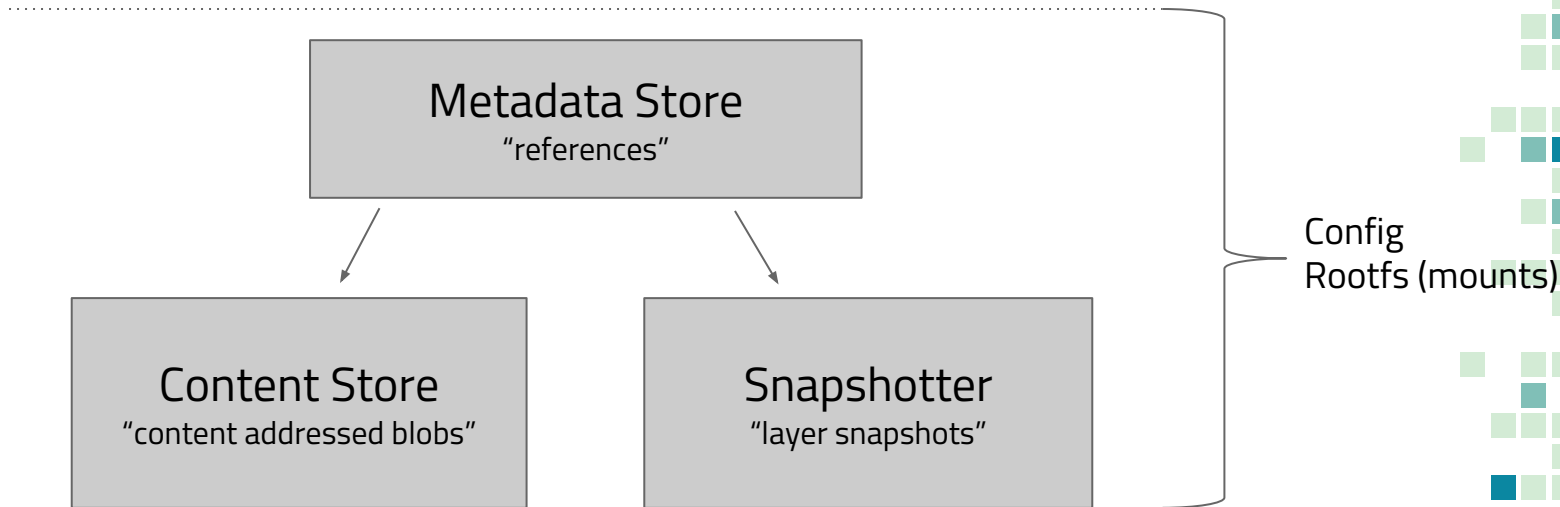
How do you build a container root filesystem?



# Docker Storage Architecture



# containerd Storage Architecture





# Snapshots

- No mounting, just returns mounts!
- Explicit active (rw) and committed (ro)
- Commands represent lifecycle
- Reference key chosen by caller (allows using content addresses)
- No tars and no diffs

## Evolved from Graph Drivers

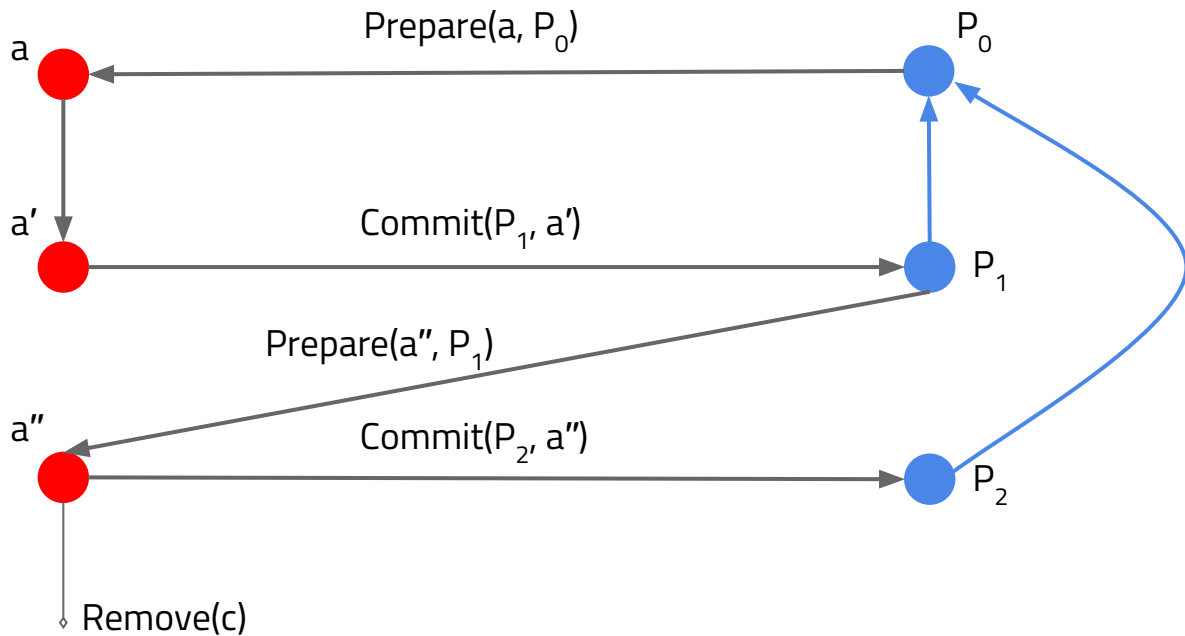
- Simple layer relationships
- Small and focused interface
- Non-opinionated string keys
- External Mount Lifecycle

```
type Snapshotter interface {  
    Stat(key string) (Info, error)  
    Mounts(key string) ([]containerd.Mount, error)  
    Prepare(key, parent string) ([]containerd.Mount, error)  
    View(key, parent string) ([]containerd.Mount, error)  
    Commit(name, key string) error  
    Remove(key string) error  
    Walk(fn func(Info) error) error  
}  
  
type Info struct {  
    Name    string // name or key of snapshot  
    Parent  string  
    Kind    Kind  
}  
  
type Kind int  
const (  
    KindView  
    KindActive  
    KindCommitted  
)
```

# Snapshot Model

**Active**

**Committed**



# Example: Investigating Root Filesystem

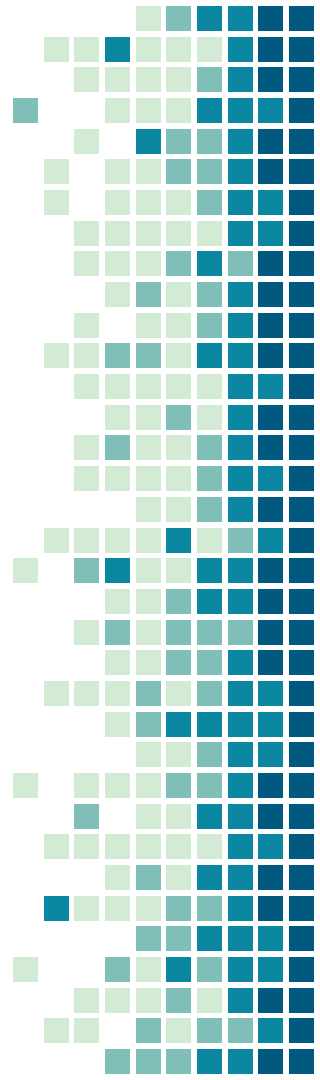
```
$ ctr snapshot ls
```

```
...
```

```
$ ctr snapshot tree
```

```
...
```

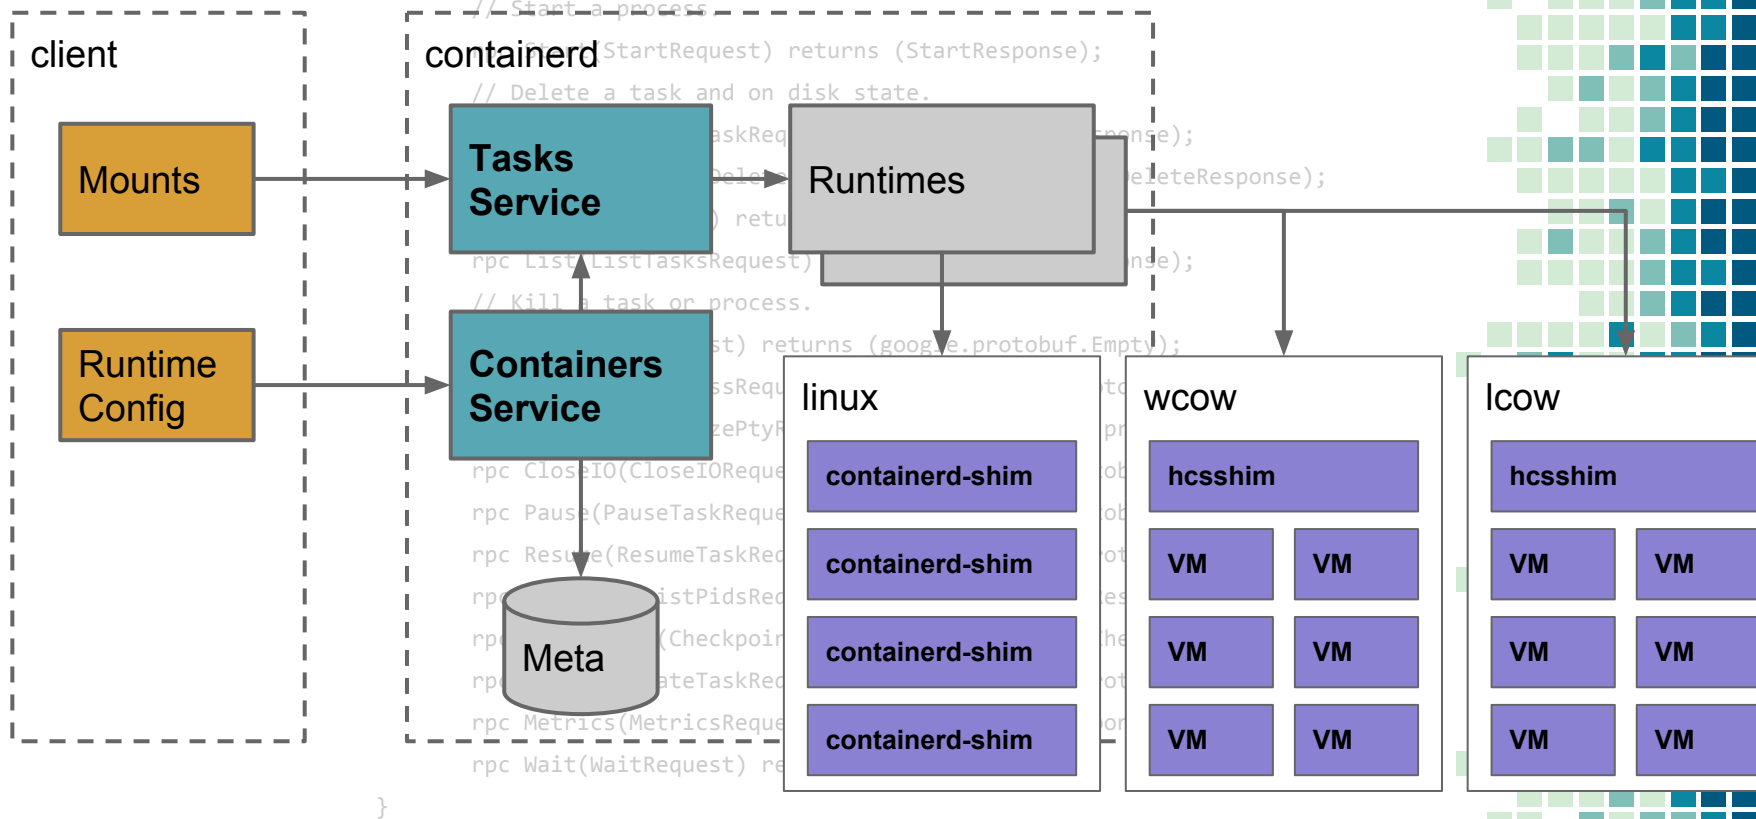
```
$ ctr snapshot mounts <target> <id>
```




Running a container



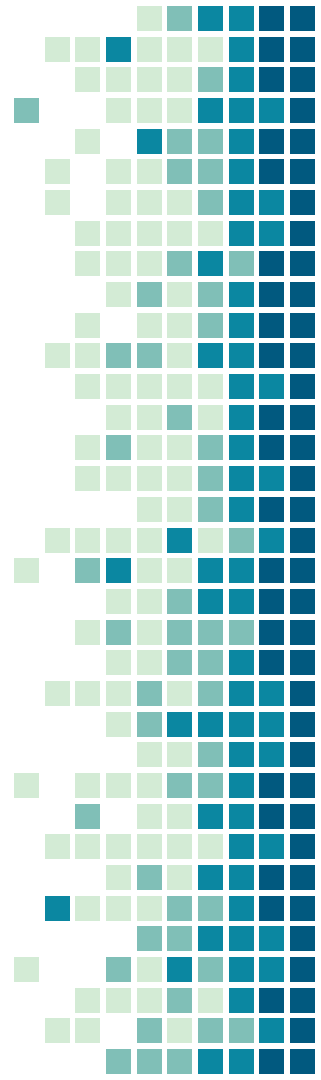

# Tasks and Runtime



# Starting a Container



Demo



# Example: Pull an Image

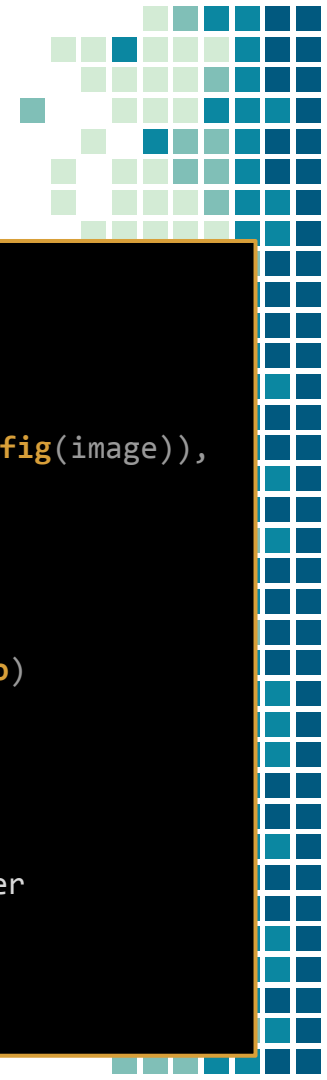
Via **ctr** client:

```
$ export \  
  CONTAINERD_NAMESPACE=example  
  
$ ctr pull \  
  docker.io/library/redis:alpine  
  
$ ctr image ls  
  
...
```

```
import (  
    "context"  
  
    "github.com/containerd/containerd"  
    "github.com/containerd/containerd/namespaces"  
)  
  
// connect to our containerd daemon  
client, err := containerd.New("/run/containerd/containerd.sock")  
defer client.Close()  
  
// set our namespace to "example":  
ctx := namespaces.WithNamespace(context.Background(), "example")  
  
// pull the alpine-based redis image from DockerHub:  
image, err := client.Pull(ctx,  
    "docker.io/library/redis:alpine",  
    containerd.WithPullUnpack)
```



# Example: Run a Container



## Via **ctr** client:

```
$ export \  
  CONTAINERD_NAMESPACE=example  
  
$ ctr run -t \  
  docker.io/library/redis:alpine \  
  redis-server  
  
$ ctr c  
...
```

```
// create our container object and config  
container, err := client.NewContainer(ctx,  
  "redis-server",  
  containerd.WithImage(image),  
  containerd.WithNewSpec(containerd.WithImageConfig(image)),  
  )  
defer container.Delete()  
  
// create a task from the container  
task, err := container.NewTask(ctx, containerd.Stdio)  
defer task.Delete(ctx)  
  
// make sure we wait before calling start  
exitStatusC, err := task.Wait(ctx)  
// call start on the task to execute the redis server  
if err := task.Start(ctx); err != nil {  
  return err  
}
```

# Example: Kill a Task



## Via **ctr** client:

```
$ export \  
  CONTAINERD_NAMESPACE=example  
  
$ ctr t kill redis-server  
  
$ ctr t ls  
...
```

```
// make sure we wait before calling start  
exitStatusC, err := task.Wait(ctx)  
  
time.Sleep(3 * time.Second)  
  
if err := task.Kill(ctx, syscall.SIGTERM); err != nil {  
    return err  
}  
  
// retrieve the process exit status from the channel  
status := <-exitStatusC  
code, exitedAt, err := status.Result()  
if err != nil {  
    return err  
}  
  
// print out the exit code from the process  
fmt.Printf("redis-server exited with status: %d\n", code)
```

# Example: Customize OCI Configuration

```
// WithHtop configures a container to monitor the host via `htop`
func WithHtop(s *specs.Spec) error {
    // make sure we are in the host pid namespace
    if err := containerd.WithHostNamespace(specs.PIDNamespace)(s); err != nil {
        return err
    }
    // make sure we set htop as our arg
    s.Process.Args = []string{"htop"}
    // make sure we have a tty set for htop
    if err := containerd.WithTTY(s); err != nil {
        return err
    }
    return nil
}
```

**With{func}** functions cleanly separate modifiers

# Release

<https://github.com/containerd/containerd/blob/master/RELEASES.md>



# Supported Components

Component	Status	Stablized Version	Links
GRPC API	Beta	1.0	<a href="#">api/</a>
Metrics API	Beta	1.0	
Go client API	Unstable	1.1 tentative	<a href="#">godoc</a>
ctr tool	Unstable	Out of scope	-



# Support Horizon

Release	Status	Start	End of Life
<a href="#">0.0</a>	End of Life	Dec 4, 2015	-
<a href="#">0.1</a>	End of Life	Mar 21, 2016	-
<a href="#">0.2</a>	Active	Apr 21, 2016	Upon 1.0 release
1.0	Next	TBD	max(TBD+1 year, release of 1.1.0)

# 1.0.0-beta.2

<https://github.com/containerd/containerd/releases/tag/v1.0.0-beta.2>



One point Ohhhhhhhhhhh!

<https://github.com/containerd/containerd/milestone/13>





# Going further with **containerd**

- **Contributing:** <https://github.com/containerd/containerd>
  - Bug fixes, adding tests, improving docs, validation
- **Using:** See the getting started documentation in the **docs** folder of the repo
- **Porting/testing:** Other architectures & OSs, stress testing (see bucketbench, containerd-stress):
  - `git clone <repo>`, `make binaries`, `sudo make install`
- **K8s CRI:** incubation project to use containerd as CRI
  - In alpha today; e2e tests, validation, contributing



# Thank You! Questions?

- **Stephen Day**
  - <https://github.com/stevvooe>
  - [stephen@docker.com](mailto:stephen@docker.com)
  - Twitter: @stevvooe

